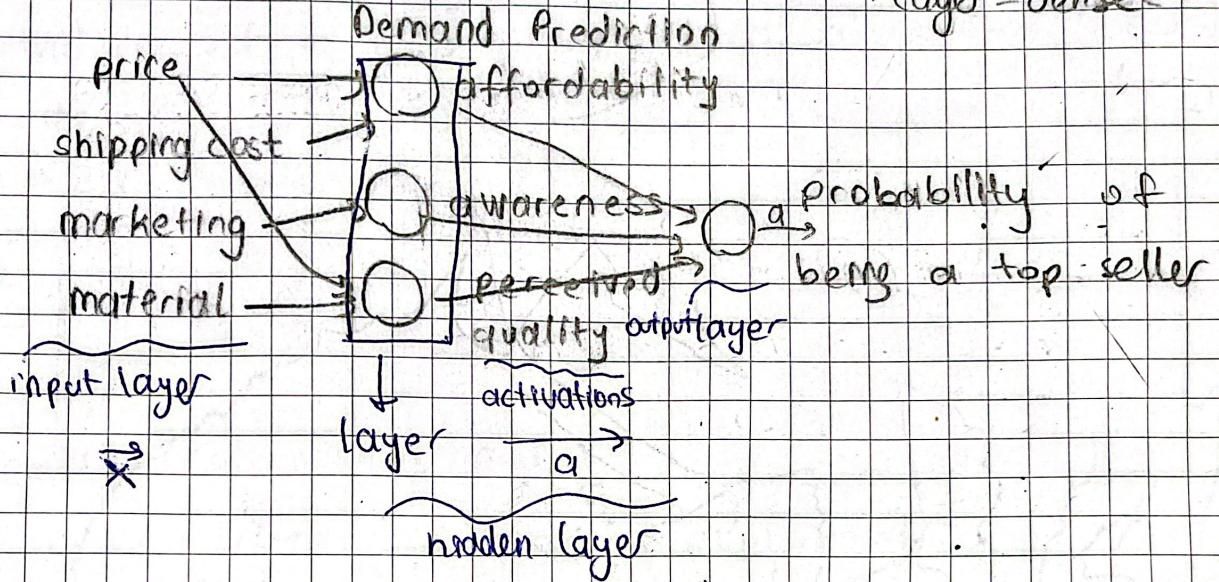
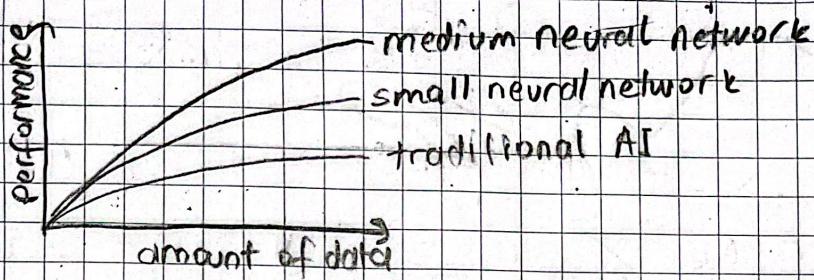


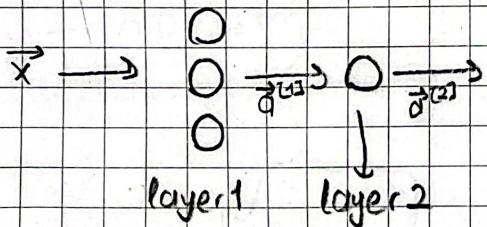
Neural Networks

Speech → images → text (NLP)



multilayer perceptron

Neural network layers



Inference: making predictions (forward propagation)

Tensorflow implementation

$x = np.array([200, 17])$ [200, 17] 1x2

$x = np.array([200, 17])$ 1D vector

```
x = np.array([[200.0, 17.0]])
```

```
layer_1 = Dense(units=3, activation="sigmoid")
```

```
a1 = layer_1(x)
```

```
layer_2 = Dense(units=1, activation="sigmoid")
```

```
a2 = layer_2(a1)
```

forward propagation

ACI AI

ANI

AGI

self-driving car Do anything a human can do

```
def dense(A_in, W, B):
```

```
Z = np.matmul(A_in, W) + B
```

```
A_out = g(Z)
```

```
return A_out
```

AT = A.T

```
Z = np.matmul(AT, W) = AT @ W
```

epochs: number of steps in gradient descent

neural network

```
model = Sequential()
```

```
Dense(...)  
Dense(...)  
])
```

binary cross entropy

```
model.compile(loss=BinaryCrossentropy())
```

```
model.fit(X, y, epochs=100)
```

Examples of Activation Functions

Linear activation function (NO activation function)

Sigmoid

ReLU (Rectified Linear Unit)

KeshinColor

Softmax regression algorithm

$$\begin{aligned}
 \times z_1 &= \vec{w}_1 \cdot \vec{x} + b_1 & a_1 &= \frac{e^{z_1}}{e^{z_1} + e^{z_2} + e^{z_3} + e^{z_4}} \\
 \circ z_2 &= \vec{w}_2 \cdot \vec{x} + b_2 & \times & \square \quad \Delta \\
 \square z_3 &= \vec{w}_3 \cdot \vec{x} + b_3 & & \\
 \Delta z_4 &= \vec{w}_4 \cdot \vec{x} + b_4 & = P(y=1|\vec{x})
 \end{aligned}$$

Adam algorithm Intuition

Adam: Adaptive Moment estimation not just one α

Additional layer types

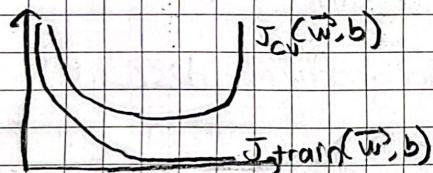
- convolutional layer

Train/test procedure for linear regression (with squared error cost)

Fit parameters by minimizing cost function $J(\vec{w}, b)$

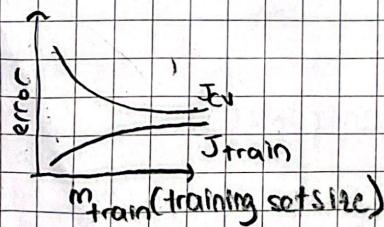
$$J(\vec{w}, b) = \frac{1}{2m_{\text{train}}} \sum_{i=1}^{m_{\text{train}}} (f_{\vec{w}, b}(\vec{x}^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m_{\text{train}}} \sum_{j=1}^n w_j^2$$

Diagnosing bias and variance



learning curves

J_{train} = training error J_{cv} = cross validation error



Debugging a learning algorithm

You've implemented regularized linear regression on housing prices but still large errors.

Get more training example \rightarrow fixes high variance

Try smaller sets of features \rightarrow high variance

Try getting additional features \rightarrow high bias

Try adding polynomial features \rightarrow fixes high bias

Try decreasing λ fix high bias

Try increasing λ fix high variance

Large neural networks are low bias machines.

Conventional
model-centric
approach

AI = Code + Data
work on this

Data-centric
approach

AI = Code + Data
work on this

Transfer learning

Decision Trees

$$P_0 = 1 - P_1$$

$$+ \frac{1}{2} \quad ①$$

$$H(P_1) = -P_1 \log_2(P_1) - P_0 \log_2(P_0)$$

$$0 \log(0) = 0$$

One-hot encoding

Tree ensemble

sampling with replacement

Random forest

XGBoost (eXtreme Gradient Boosting)

Decision trees vs Neural Networks

KashinColor